

Lab 1 Prometheus

1-what is different http status code and explain meaning of each of them?

1. Informational responses (100 – 199)
2. Successful responses (200 – 299)
3. Redirection messages (300 – 399)
4. Client error responses (400 – 499)
5. Server error responses (500 – 599)

Information responses

100 Continue

The client should continue the request or ignore the response if the request is already finished.

101 Switching Protocols

This code is sent in response to an Upgrade request header from the client and indicates the protocol the server is switching to.

102 Processing

This code indicates that the server has received and is processing the request, but no response is available yet.

103 Early Hints Experimental

This status code is primarily intended to be used with the Link header, letting the user agent start preloading resources while the server prepares a response.

Successful responses

200 OK

The request succeeded. The result meaning of "success" depends on the HTTP method:
GET - HEAD - POST - PUT - TRACE

201 Created

The request succeeded, and a new resource was created as a result.

This is typically the response sent after POST requests, or some PUT requests.

202 Accepted

The request has been received but not yet acted upon.

203 Non-Authoritative Information

This response code means the returned metadata is not exactly the same as is available from the origin server,

but is collected from a local or a third-party copy.

This is mostly used for mirrors or backups of another resource. Except for that specific case, the 200 OK response is preferred to this status.

204 No Content

There is no content to send for this request,

but the headers may be useful. The user agent may update its cached headers for this resource with the new ones.

205 Reset Content

Tells the user agent to reset the document which sent this request.

206 Partial Content

This response code is used when the Range header is sent from the client to request only part of a resource.

207 Multi-Status (WebDAV)

Conveys information about multiple resources, for situations where multiple status codes might be appropriate.

208 Already Reported (WebDAV)

Used inside a < dav:propstat> response element to avoid repeatedly enumerating the internal members of multiple bindings to the same collection.

226 IM Used (HTTP Delta encoding)

The server has fulfilled a GET request for the resource,

and the response is a representation of the result of one or more instance-manipulations applied to the current instance.

Redirection messages

300 Multiple Choices

The request has more than one possible response.

The user agent or user should choose one of them.

(There is no standardized way of choosing one of the responses, but HTML links to the possibilities are recommended so the user can pick.)

301 Moved Permanently

The URL of the requested resource has been changed permanently. The new URL is given in the response.

302 Found

This response code means that the URI of requested resource has been changed temporarily.

Further changes in the URI might be made in the future. Therefore, this same URI should be used by the client in future requests.

303 See Other

The server sent this response to direct the client to get the requested resource at another URI with a GET request.

304 Not Modified

This is used for caching purposes.

It tells the client that the response has not been modified, so the client can continue to use the same cached version of the response.

305 Use Proxy Deprecated

Defined in a previous version of the HTTP specification to indicate that a requested response must be accessed by a proxy. It has been deprecated due to security concerns regarding in-band configuration of a proxy.

306 unused

This response code is no longer used; it is just reserved. It was used in a previous version of the HTTP/1.1 specification.

307 Temporary Redirect

The server sends this response to direct the client to get the requested resource at another URI with the same method that was used in the prior request.

308 Permanent Redirect

This means that the resource is now permanently located at another URI, specified by the Location: HTTP Response header.

Client error responses

400 Bad Request

The server cannot or will not process the request due to something that is perceived to be a client error (e.g., malformed request syntax, invalid request message framing, or deceptive request routing).

401 Unauthorized

Although the HTTP standard specifies "unauthorized", semantically this response means "unauthenticated". That is, the client must authenticate itself to get the requested response.

402 Payment Required Experimental

This response code is reserved for future use. The initial aim for creating this code was using it for digital payment systems, however this status code is used very rarely and no standard convention exists.

403 Forbidden

The client does not have access rights to the content; that is, it is unauthorized, so the server is refusing to give the requested resource. Unlike 401 Unauthorized, the client's identity is known to the server.

404 Not Found

The server cannot find the requested resource. In the browser, this means the URL is not recognized. In an API, this can also mean that the endpoint is valid but the resource itself does not exist. Servers may also send this response instead of 403 Forbidden to hide the existence of a resource from an unauthorized client. This response code is probably the most well known due to its frequent occurrence on the web.

405 Method Not Allowed

The request method is known by the server but is not supported by the target resource. For example, an API may not allow calling DELETE to remove a resource.

406 Not Acceptable

This response is sent when the web server, after performing server-driven content negotiation, doesn't find any content that conforms to the criteria given by the user agent.

407 Proxy Authentication Required

This is similar to 401 Unauthorized but authentication is needed to be done by a proxy.

408 Request Timeout

This response is sent on an idle connection by some servers, even without any previous request by the client. It means that the server would like to shut down this unused connection. This response is used much more since some browsers, like Chrome, Firefox 27+, or IE9, use HTTP pre-connection mechanisms to speed up surfing. Also note that some servers merely shut down the connection without sending this message.

409 Conflict

This response is sent when a request conflicts with the current state of the server.

410 Gone

This response is sent when the requested content has been permanently deleted from server, with no forwarding address. Clients are expected to remove their caches and links to

the resource. The HTTP specification intends this status code to be used for "limited-time, promotional services". APIs should not feel compelled to indicate resources that have been deleted with this status code.

411 Length Required

Server rejected the request because the Content-Length header field is not defined and the server requires it.

412 Precondition Failed

The client has indicated preconditions in its headers which the server does not meet.

413 Payload Too Large

Request entity is larger than limits defined by server. The server might close the connection or return an Retry-After header field.

414 URI Too Long

The URI requested by the client is longer than the server is willing to interpret.

415 Unsupported Media Type

The media format of the requested data is not supported by the server, so the server is rejecting the request.

416 Range Not Satisfiable

The range specified by the Range header field in the request cannot be fulfilled.

It's possible that the range is outside the size of the target URI's data.

417 Expectation Failed

This response code means the expectation indicated by the Expect request header field cannot be met by the server.

418 I'm a teapot

The server refuses the attempt to brew coffee with a teapot.

421 Misdirected Request

The request was directed at a server that is not able to produce a response.

422 Unprocessable Entity (WebDAV)

The request was well-formed but was unable to be followed due to semantic errors.

423 Locked (WebDAV)

The resource that is being accessed is locked.

424 Failed Dependency (WebDAV)

The request failed due to failure of a previous request.

425 Too Early Experimental

Indicates that the server is unwilling to risk processing a request that might be replayed.

426 Upgrade Required

The server refuses to perform the request using the current protocol but might be willing to do so after the client upgrades to a different protocol.

The server sends an Upgrade header in a 426 response to indicate the required protocol(s).

428 Precondition Required

The origin server requires the request to be conditional. This response is intended to prevent the 'lost update' problem,

where a client GETs a resource's state, modifies it and PUTs it back to the server, when meanwhile a third party has modified the state on the server,

leading to a conflict.

429 Too Many Requests

The user has sent too many requests in a given amount of time ("rate limiting").

431 Request Header Fields Too Large

The server is unwilling to process the request because its header fields are too large.

The request may be resubmitted after reducing the size of the request header fields.

451 Unavailable For Legal Reasons

The user agent requested a resource that cannot legally be provided, such as a web page censored by a government.

[Server error responses](#)

500 Internal Server Error

The server has encountered a situation it does not know how to handle.

501 Not Implemented

The request method is not supported by the server and cannot be handled.

The only methods that servers are required to support (and therefore that must not return this code) are GET and HEAD.

502 Bad Gateway

This error response means that the server, while working as a gateway to get a response needed to handle the request, got an invalid response.

503 Service Unavailable

The server is not ready to handle the request.

Common causes are a server that is down for maintenance or that is overloaded.

504 Gateway Timeout

This error response is given when the server is acting as a gateway and cannot get a response in time.

505 HTTP Version Not Supported

The HTTP version used in the request is not supported by the server.

506 Variant Also Negotiates

The server has an internal configuration error: the chosen variant resource is configured to engage in transparent content negotiation itself,

and is therefore not a proper end point in the negotiation process.

507 Insufficient Storage

The method could not be performed on the resource because the server is unable to store the representation needed to successfully complete the request.

508 Loop Detected

The server detected an infinite loop while processing the request.

510 Not Extended

Further extensions to the request are required for the server to fulfill it.

511 Network Authentication Required

Indicates that the client needs to authenticate to gain network access.

2- What database is used by Prometheus?

Prometheus stores the metrics it scrapes on a local disk in a time-series database.

(Prometheus includes a local on-disk time series database, but also optionally integrates with remote storage systems).

3- What is the difference between different metrics types (counter , gauge , histogram)?

Counter

A counter is a cumulative metric that represents a single monotonically increasing counter whose value can only increase or be reset to zero on restart.

For example, you can use a counter to represent the number of requests served, tasks completed, or errors.

Gauge

A gauge is a metric that represents a single numerical value that can arbitrarily go up and down

Histogram

A histogram samples observations (usually things like request durations or response sizes) and counts them in configurable buckets.

It also provides a sum of all observed values.

4- Install Prometheus on your localhost or on server in any cloud provider

```
● prometheus.service - Prometheus
   Loaded: loaded (/etc/systemd/system/prometheus.service; enabled; vendor preset: ena
   Active: active (running) since Sat 2023-02-18 00:46:21 UTC; 37min ago
     Main PID: 29532 (prometheus)
        Tasks: 8 (limit: 1143)
       Memory: 47.8M
          CPU: 3.312s
      CGroup: /system.slice/prometheus.service
              └─29532 /usr/local/bin/prometheus --config.file /etc/prometheus.conf

Feb 18 00:46:24 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T00:46:24.123Z
Feb 18 00:46:24 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T00:46:24.123Z
Feb 18 00:46:24 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T00:46:24.123Z
Feb 18 00:46:24 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T00:46:24.123Z
Feb 18 00:46:24 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T00:46:24.123Z
Feb 18 00:46:24 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T00:46:24.123Z
Feb 18 01:00:10 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T01:00:10.123Z
Feb 18 01:00:10 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T01:00:10.123Z
Feb 18 01:00:10 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T01:00:10.123Z
Feb 18 01:00:10 ip-172-31-42-109 prometheus[29532]: ts=2023-02-18T01:00:10.123Z
lines 1-20/20 (END)
```

5- Add any new target to prometheus.yml file and apply any query on it using promql language

The screenshot shows the Prometheus web interface at localhost:9090. The query bar contains the PromQL query: `container_cpu_usage_seconds_total{image="gcr.io/cadvisor/cadvisor"}`. The interface includes checkboxes for "Use local time", "Enable query history", "Enable autocomplete", "Enable highlighting", and "Enable linter". Below the query bar, there are tabs for "Table" and "Graph". The "Table" tab is selected, showing a single data point with a value of 341.40443. The data point includes labels for `cpu`, `id`, `instance`, and `name`. The interface also has an "Execute" button and a "Remove Panel" link.

container_cpu_usage_seconds_total
341.40443